

# BigSem: Big Data Analytics for Semantic Data Tutorial

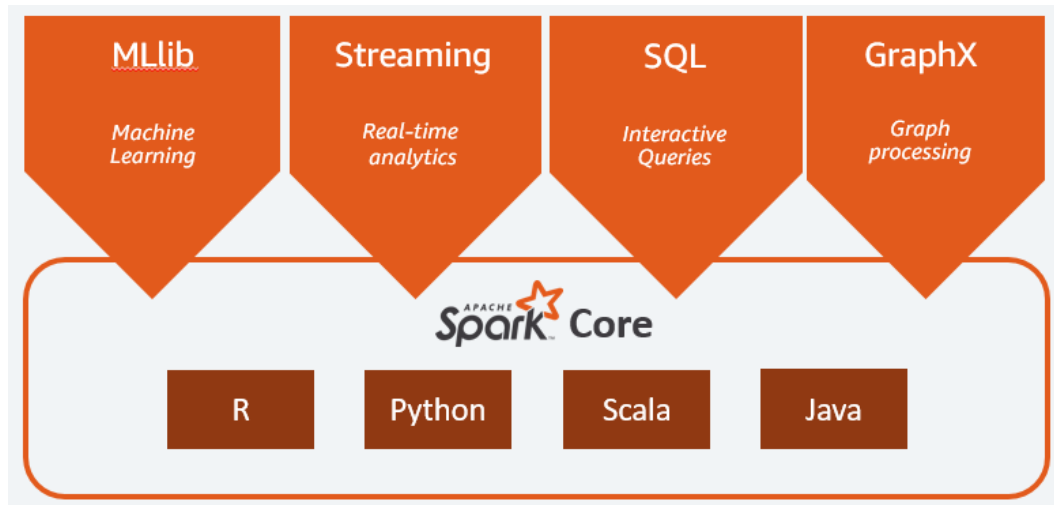
## Module 3: Semantic data analytic engines and frameworks

Chelmis Charalampos, Bedirhan Gergin  
University at Albany, SUNY

*ISWC 2024*


# Apache Spark

- Apache Spark is an open-source cluster computing system that provides high-level API in Java, Scala, Python.



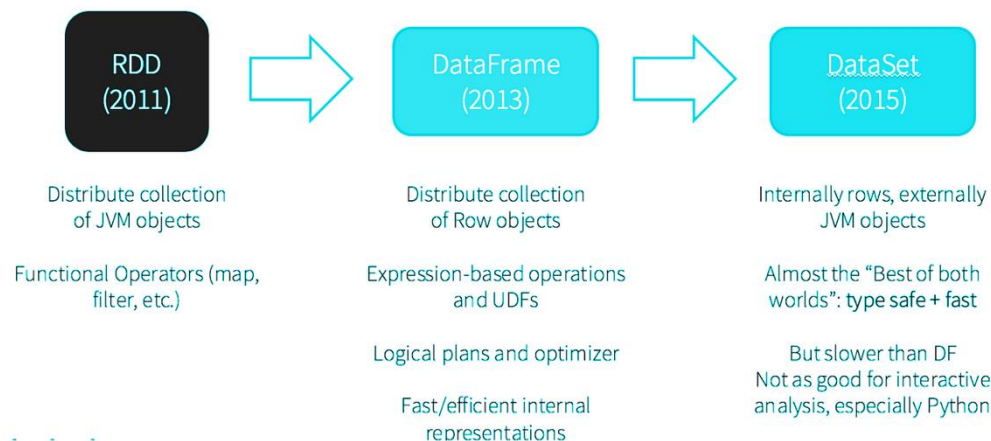
# Apache Spark: Architecture

- **Driver Program:** The Spark Context manages job execution and distributes tasks to the worker nodes.
- **Cluster Manager:** Allocates resources across the cluster. It communicates with both the driver program and the worker nodes to manage resource allocation.
- **Worker Nodes:** where the actual execution of tasks happens. Each worker node contains executors, which are responsible for running the tasks.




# Apache Spark: RDD

- An RDD (Resilient Distributed Dataset) in Apache Spark is a fundamental data structure that represents an immutable, distributed collection of objects that can be processed in parallel across a cluster.



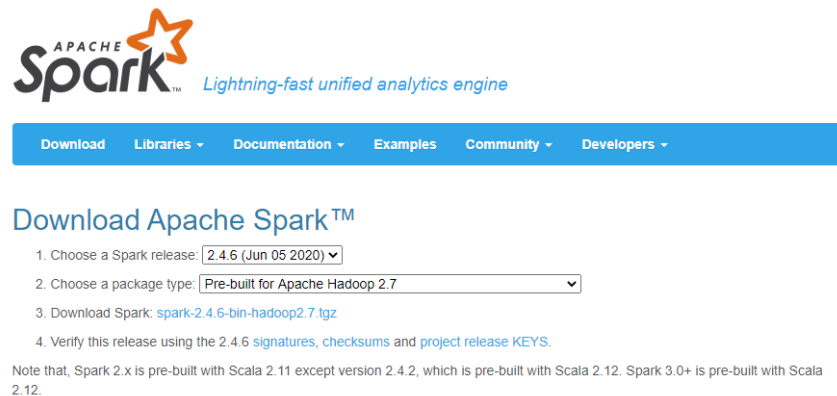
# Apache Spark in Python: PySpark

- PySpark is the Python API for Apache Spark. It enables you to perform real-time, large-scale data processing in a distributed environment using Python.



# Installing Spark

- Head over to the Spark homepage.
- Select the Spark release and package type as following and download the .tgz file.
- Configure Environment Variable for Apache Spark and Python



```
export SPARK_HOME="/Downloads/spark"
export PATH=$SPARK_HOME/bin:$PATH
export PYSARK_PYTHON=python3
```

# Install PySpark

- Install PySpark using pip

```
pip install findspark  
pip install pyspark
```

Welcome to

```
  _ _ _ _ _  
 / _ \ _ _ _ \ _ /  
 \ V _ V _ \ _ /  
 / _ / . _ \ , / / \ \ version 2.4.  
 / _ \
```

# PySpark: SparkSession

- **config:** This allows users to specify additional Spark configuration settings to customize the application further.
- **spark.executor.memory:** This parameter specifies the amount of memory allocated per executor process, such as 2g for 2 gigabytes.
- **spark.executor.cores:** This defines the number of CPU cores allocated to each executor, impacting the parallelism and speed of the application.
- **spark.driver.memory:** This indicates the amount of memory reserved for the driver process, with a common setting being 1g.

```
import findspark
findspark.init()
from pyspark.sql import SparkSession

# Initialize SparkSession with additional configurations
spark = SparkSession.builder \
    .appName('PySpark Beginner Tutorial') \
    .master('local[*]') \
    .config('spark.executor.memory', '2g') \
    .config('spark.executor.cores', '2') \
    .config('spark.driver.memory', '1g') \
    .getOrCreate()
```



# PySpark: Creating Dataframe

- Creating spark dataframe manually.

```
data = [('John', 28, 'M'), ('Anna', 23, 'F'), ('Mike', 35, 'M'), ('Sara', 31, 'F')]
columns = ['Name', 'Age', 'Sex']

# Create a DataFrame with additional columns
df = spark.createDataFrame(data, columns)

# Show DataFrame
df.show()
```

```
+---+---+---+
|Name|Age|Sex|
+---+---+---+
|John| 28|  M|
|Anna| 23|  F|
|Mike| 35|  M|
|Sara| 31|  F|
+---+---+---+
```

# PySpark: Dataframe Operations

- Here are some common DataFrame operations such as checking schema, selecting columns, filtering rows, and computing basic statistics.

```
# Check the schema of the DataFrame
df.printSchema()

# Select the 'Name' column
df.select('Name').show()

# Filter rows where age > 30
df.filter(df.Age > 30).show()

# Compute basic statistics
df.describe().show()
```

```
root
 |-- Name: string (nullable = true)
 |-- Age: long (nullable = true)
 |-- Sex: string (nullable = true)


+----+
|Name|
+----+
|John|
|Anna|
|Mike|
|Sara|
+----+

+---+---+---+
|Name|Age|Sex|
+---+---+---+
|Mike| 35|  M|
|Sara| 31|  F|
+---+---+---+

+---+---+---+
|summary| Age|
+---+---+---+
|  count|   4|
|   mean|29.25|
| stddev|4.573474244670772|
|   min|  23|
|   max|  35|
+---+---+---+
```

# SANSA Stack

- Scalable Semantic Analytic Stack (SANSA) framework is an open-source distributed data flow engine which allows scalable analysis of large-scale RDF datasets.



# SANSA Stack – RDF Processing Layer

- SANSA provide mechanism of reading RDF model in the format of RDD/DataFrame/Dataset of triples.

Listing 1. Triple reader example.

```
1 import net.sansa_stack.rdf.spark.io._
2 import org.apache.jena.riot.Lang
3
4 val input = "hdfs://namenode:8020/data/rdf.nt"
5 val lang = Lang.NTRIPLES
6
7 val triples = spark.rdf(lang)(input)
8
9 triples.take(5).foreach(println(_))
```

Listing 2. Triple writer example.

```
1 import net.sansa_stack.rdf.spark.io._
2 import org.apache.jena.riot.Lang
3
4 val input = "hdfs://namenode:8020/data/rdf.nt"
5 val lang = Lang.NTRIPLES
6
7 val triples = spark.rdf(lang)(input)
8
9 triples.saveAsNTriplesFile(output)
```

# SANSA Stack – Querying Layer

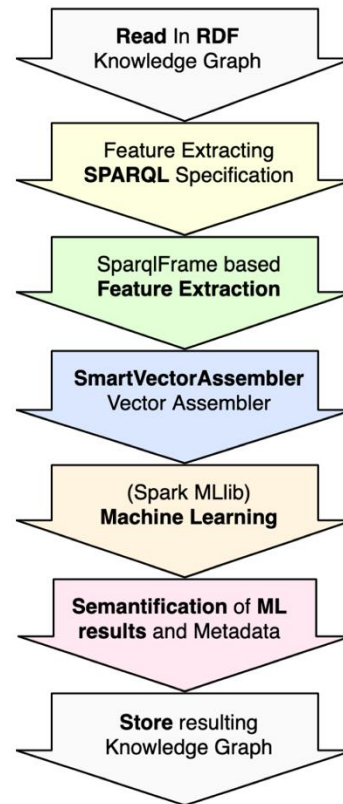
- The default approach for querying RDF data in SANSA is based on SPARQL-to-SQL. It uses a flexible triple-based partitioning strategy on top of RDF (such as predicate tables with sub partitioning by data types).

Listing 8. Sparklify example.

```
1 import org.apache.jena.riot.Lang
2 import net.sansa_stack.rdf.spark.io._
3 import net.sansa_stack.query.spark.query._
4
5 val input = "hdfs://namenode:8020/data/rdf.nt"
6 val lang = Lang.NTRIPLES
7
8 val triples = spark.rdf(lang)(input)
9 val sparqlQuery = """SELECT ?s ?p ?o
10                      WHERE {?s ?p ?o }
11                      LIMIT 10"""
12 val result = triples.sparql(sparqlQuery)
13 z.show(result)
```

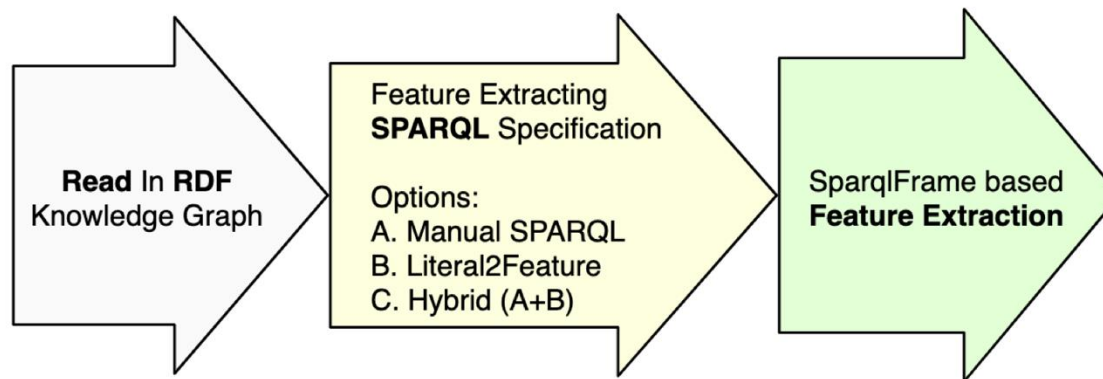
# SANSA Stack – ML Layer (DistRDF2ML)

- Introduces software modules that transform large-scale RDF data into ML-ready fixed-length numeric feature vectors




# SANSA Stack – DistRDF2ML (SparqlFrame)

- **Manual SPARQL creation:** A knowledge graph expert manually crafts the SPARQL query to extract relevant features.
- **Literal2Feature:** automatically generates a SPARQL query by deep traversing the RDF graph and extracting literals.
- **Hybrid approach:** Literal2Feature generates a query, which is then manually refined, balancing automation with manual control to create a clean and focused SPARQL query.



# SANSA Stack – DistRDF2ML (SmartVectorAssembler)


- **SparqlFrame**: Convert SPARQL query results into Spark DataFrames.
- **SmartVectorAssembler**: Features are converted into numeric representations based on their type (e.g., Word2Vec for strings, indices for categorical lists, and datetime transformations for timestamps).





# SparkKG-ML

- A Library to Facilitate end-to-end Large-scale Machine Learning over Knowledge Graphs in Python.



# SparkKG-ML

**Data Acquisition:** Transform RDF data into the tabular format that Spark can process.

**Feature Eng.:** Gathers feature characteristics and a collapsed DataFrame is created.

**Feature Selection:** Focuses on identifying and retaining attributes while discarding redundant ones.

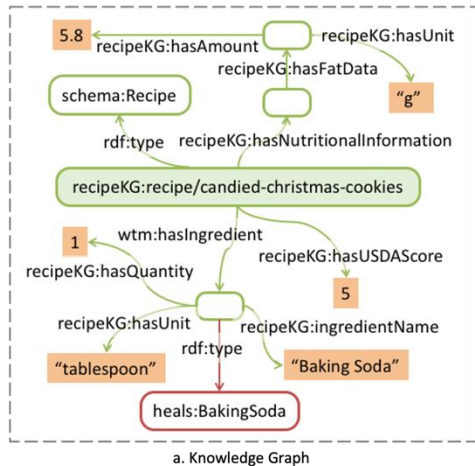
**Data Augmentation:** Enables augmenting a given KG with data from public KGs, allowing the extraction of additional features from these KGs.

**Vectorization:** Produces a ML-ready DataFrame by transforming all features according to feature type into numeric representations.

**Semantification:** Transform ML results into RDF data.

# SparkKG-ML: Data Acquisition

- Transform RDF data into the tabular format that Spark can process.



b. SPARQL Query

```

PREFIX schema: <https://schema.org/>
PREFIX recipeKG: <http://purl.org/recipekg/>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
PREFIX food: <http://purl.org/heals/food/>
SELECT ?recipe ?ingredientName ?fat
WHERE { ?recipe a schema:Recipe.
        ?recipe food:hasIngredient ?ingredient.
        ?ingredient recipeKG:ingredientName ?ingredientName.
        ?recipe recipeKG:hasNutritionalInformation ?a.
        ?a recipeKG:hasFatData ?b.
        ?b recipeKG:hasAmount ?fat. }
    
```

c. Query Result

recipe	ingredientName	fat
1	recipeKG:recipe/candied-christmas-cookies	"all purpose flour" "5.8"^^xsd:float
2	recipeKG:recipe/candied-christmas-cookies	"baking soda" "5.8"^^xsd:float
3	recipeKG:recipe/candied-christmas-cookies	"bourbon" "5.8"^^xsd:float
4	recipeKG:recipe/candied-christmas-cookies	"brown sugar" "5.8"^^xsd:float
5	recipeKG:recipe/candied-christmas-cookies	"butter" "5.8"^^xsd:float
6	recipeKG:recipe/peanut-butter-tandy-bars	"egg" "9.5"^^xsd:float
7	recipeKG:recipe/peanut-butter-tandy-bars	"butter" "9.5"^^xsd:float
8	recipeKG:recipe/peanut-butter-tandy-bars	"chocolate" "9.5"^^xsd:float
9	recipeKG:recipe/peanut-butter-tandy-bars	"baking powder" "9.5"^^xsd:float
10	recipeKG:recipe/the-best-oatmeal-cookies	"cinnamon" "7.6"^^xsd:float

```

# Import the required module
from sparkkgml.data_acquisition import DataAcquisition

# Create an instance of DataAcquisition
dataAcquisitionObject=DataAcquisition()

# Specify the SPARQL endpoint and query
endpoint = "https://recipekg.arcc.albany.edu/RecipeKG"
query = "" ... ""

# Retrieve the data as a Spark DataFrame
spark_df = dataAcquisitionObject.getDataFrame(endpoint=
    endpoint, query=query)
    
```

recipe	ingredient	fat
candied-chri...	flour	5.8
candied-chri...	baking soda	5.8
candied-chri...	bourbon	5.8
candied-chri...	brown sugar	5.8
candied-chri...	butter	5.8
peanut-butte...	egg	9.5
peanut-butte...	butter	9.5
peanut-butte...	chocolate	9.5
peanut-butte...	baking powder	9.5
the-best-oat...	cinnamon	7.6

# SparkKG-ML: Feature Engineering

- Gathers feature characteristics and a collapsed DataFrame is created.
  - **datatype**: The data type of the feature column.
  - **numberDistinctValues**: The number of distinct values in the feature column.
  - **isListOfEntries**: Flag indicating if the feature is a list of entries.
  - **isCategorical**: The ratio of distinct values and overall dataset size
  - **featureType**: Combine features based on whether they consist of a list or a single value, categorical or non-categorical, and data type.

```
# Import the required module
from sparkkgml.feature_engineering import FeatureEngineering
from sparkkgml.vectorization import Vectorization

# Create an instance of FeatureEngineering
featureEngineeringObject=FeatureEngineering()

# Call the getFeatures function
df2,features=featureEngineeringObject.getFeatures(spark_df)

# Create an instance of Vectorization
vectorizationObject=Vectorization()

# Call vectorize function, digitize all the columns
digitized_df=vectorizationObject.vectorize(df2,features)
```

recipe	ingredients	fat
candied-chri...	[baking soda, egg, b...	5.8
peanut-butte...	[egg, butter, chocol...	9.5
best-oatmeal...	[cinnamon, egg, suga...	7.6
alfredo-blue...	[salt, egg, pasta, b...	5.2
millie-pasqu...	[lemon, flour, salt,...	4.3

# SparkKG-ML: Vectorization

- Gathers feature characteristics and a collapsed DataFrame is created.
  - **Single Categorical String:** indexing or hashing
  - **List of Categorical Strings:** explodes the list and applies string indexing or hashing
  - **Single Non-Categorical String:** Word2Vec (optional stop word removal)
  - **List of Non-Categorical Strings:** the list elements are combined, tokenized, stop words are removed. Embeddings are then calculated using Word2Vec.
  - **Numeric Type :** (i.e., integer, long, float, double)
  - **Boolean Type:** cast to integers (0 or 1).

entity	features	label
candied-chri...	[0.01, 1.34, 6...	5
peanut-butte...	[0.03, 6.34, 7...	6
best-oatmeal...	[0.01, 8.34, 3...	1
alfredo-blue...	[0.05, 3.34, 2...	5
millie-pasqu...	[0.61, 9.34, 1...	4